

Benchmark Guide for mTCP with DSA

Benchmark Guide for mTCP with DSA

1. Hardware Environmental Preparation

2. Software Environmental Preparation

- Configure demo machine(spr machine run demo)
 - Modify The Configuration File
 - Configure DSA (Reconfiguration is required when restart)
 - Build Test Program "Epping"
- Configure pktgen machine(machine run pktgen to send packets)
- Run Demo

4. Start Test

- Sample Run
- benchmark with CPU
- benchmark with DSA

5.Measurement

- PPS Measurement
- CPU Usage
- CPU Usage with DSA
- Bandwidth Measurement
 - 1. Peak value measurement
 - 2. Perf measurement

1. Hardware Environmental Preparation

- Two machines directly connected via 100G network card.
- One of the two must be a SPR machine, and this SPR machine kernel(v5.13) support DSA
 - To compile the kernel, you can refer to this document:
"**kernel_build_and_iax_enable_v1.0.pdf**" (The steps for compiling the kernel are universal)
 - Read the **README.md** in the **dsa_userlib** directory to install the necessary libraries.
Tips:
 - when you clone : <https://github.com/intel/idxd-config.git> , you should switch the version, do not use the latest version: you can
git checkout 6bd68e68;
 - Modify config file : /etc/ld.so.conf , add a new line : ".", and then run command
"ldconfig", Configure the current directory./ to the library search path .
 - Then follow the readme to continue.
- SPR machine BIOS Setting

BIOS Setting

EDKII Menu > Socket Configuration > Uncore Configuration > Uncore Dfx Configuration:
Cache entries for non-atomics = 120

EDKII Menu > Socket Configuration > Uncore Configuration > Uncore Dfx Configuration:
Cache entries for atomics = 8

EDKII Menu > Socket Configuration > Uncore Configuration > Uncore Dfx Configuration:
CTAG entry avail mask = 255

EDKII Menu-> "Socket Configuration" -> "IIO Configuration" -> "Intel VT for directed IO
(VT -d)" → Intel VT for directed IO → Enable

EDKII Menu-> "Socket Configuration" -> "IIO Configuration" -> "PCI ENQCMD/ENQCMTDS"
→ Enable**

2. Software Environmental Preparation

Configure demo machine(spr machine run demo)

1. Record Information

View and record the information of the directly connected network card

Suppose, you are directly connected as follows:

192.168.1.1 (pktgen machine) NIC: ens5f1 <-----NIC direct connection-----> **192.168.1.2 (demo machine) NIC: ens66(0000:29:00.0)**

You should record these information, for example:

```
[root@SPR06 mtcp_merge_virtq]# ethtool -i ens66
driver: ice
version: 5.13.0+
firmware-version: 2.40 0x80007063 1.2898.0
expansion-rom-version:
bus-info: 0000:29:00.0
supports-statistics: yes
supports-test: yes
supports-eprom-access: yes
supports-register-dump: yes
supports-priv-flags: yes
```

```
[root@SPR06 mtcp_merge_virtq]# cat /sys/class/net/ens66/address
40:a6:b7:67:19:f0
```

or enter "ifconfig":

```
ens66: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
inet 192.168.111.234 netmask 255.255.255.0 broadcast 192.168.111.255
inet6 fe80::afe:2bf7:ac5e:33a prefixlen 64 scopeid 0x20<link>
ether 40:a6:b7:67:19:f0 txqueuelen 1000 (Ethernet)
RX packets 0 bytes 0 (0.0 B)
RX errors 0 dropped 0 overruns 0 frame 0
TX packets 45 bytes 4753 (4.6 KiB)
TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

NIC: ens66

ens66 bus-info:0000:29:00.0

ens66 mac address: 40:A6:B7:67:19:F0

then, uninstall an existing drive

```
ifconfig ens66 down
```

2. Download Source Code

```
# repo: https://github.com/intel-collab/libraries.performance.accelerators.data-
streaming.dsa-share-memory-framework
# Source Code Main Directory Structure
  dsa_userlib # DSA library
    ├──config_dsa # config file is used to dsa
    ├──include # dsa library headfile
    ├──src # dsa library source code
    ├──test # dsa test program, you can read dsa readme.md to learn how to run it
  mtcp_merge_virtq # mtcp source folder
    ├──apps
      ├──example # mtcp app example folder
      │   ├──epping.c # mtcp app demo
      │   ├──epping.conf # mtcp app demo config file
    ├──mtcp # mtcp lib source code
    ├──dpdk # dpdk lib souce code
    ...

git clone git@github.com:intel-collab/libraries.performance.accelerators.data-
streaming.dsa-share-memory-framework.git
# Get into source code path
cd /dsa-accelerated-mtcp/mtcp_merge_virtq
# Run prepare shell (only need run once)
./compile_prepare.sh
```

Bind NIC with igb_uio Driver (Need to reconfigure after reboot)

```
# https://doc.dpdk.org/dts/gsg/usr_guide/igb_uio.html
1. Igb_uio driver can build by dpdk
2. Bind nic with igb_uio driver. This network card is directly connected to
another machine.

TIPS:
1. If you need to use a latest version of dpdk to compile IGB_ UIO driver,
please refer to: https://doc.dpdk.org/dts/gsg/usr_guide/igb_uio.html
2. You can also use default dpdk(Not the latest version) follow the steps below
(after unzip /dsa-accelerated-mtcp/mtcp_merge_virtq/dpdk/x86_64-native-linuxapp-
gcc.tar.gz):
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/
export RTE_SDK=`echo $PWD`/dpdk
export RTE_TARGET=x86_64-native-linuxapp-gcc
./setup_mtcp_dpdk_env.sh $RTE_SDK
1. build x86_64-native-linuxapp-gcc (for this dpdk version , enter 38 to build
it,
If there is an error in this step, you can download a separate dpdk kmod repo
and compile it)
```

```
[root@SPR06 mtcp_merge_virtq]# ./setup_mtcp_dpdk_env.sh $RTE_SDK
```

```
[36] x86_64-native-freebsd-gcc
[37] x86_64-native-linuxapp-clang
[38] x86_64-native-linuxapp-gcc
[39] x86_64-native-linuxapp-icc
[40] x86_64-native-linux-clang
[41] x86_64-native-linux-gcc
[42] x86_64-native-linux-icc
[43] x86_x32-native-linuxapp-gcc
[44] x86_x32-native-linux-gcc
```

Step 2: Setup linux environment

```
[45] Insert IGB UIO module
[46] Insert VFIO module
[47] Insert KNI module
[48] Setup hugepage mappings for non-NUMA systems
[49] Setup hugepage mappings for NUMA systems
[50] Display current Ethernet/Baseband/Crypto device settings
[51] Bind Ethernet/Baseband/Crypto device to IGB UIO module
[52] Bind Ethernet/Baseband/Crypto device to VFIO module
[53] Setup VFIO permissions
```

Step 3: Run test application for linux environment

```
[54] Run test application ($RTE_TARGET/app/test)
[55] Run testpmd application in interactive mode ($RTE_TARGET/app/testpmd)
```

Step 4: Other tools

```
[56] List hugepage info from /proc/meminfo
```

Step 5: Uninstall and system cleanup

```
[57] Unbind devices from IGB UIO or VFIO driver
[58] Remove IGB UIO module
[59] Remove VFIO module
[60] Remove KNI module
[61] Remove hugepage mappings
```

```
[62] Exit Script
```

```
Option: 38
```

```
export RTE_SDK=/root/dpdk-stable-18.11.11
```

```
== Build drivers/raw/dpaa2_qdma
== Build drivers/raw/ufpga_rawdev
== Build app
== Build app/test-pmd
== Build app/proc-info
== Build app/pdump
== Build app/test-bbdev
== Build app/test-crypto-perf
== Build app/test-eventdev
Build complete [x86_64-native-linuxapp-gcc]
Installation cannot run with T defined and DESTDIR undefined
```

RTE_TARGET exported as x86_64-native-linuxapp-gcc

```
Press enter to continue ...
```

2. Insert IGB UIO module (for this dpdk version , enter 45 to build it)

```
-----
Step 2: Setup linux environment
-----
[45] Insert IGB UIO module
[46] Insert VFIO module
[47] Insert KNI module
[48] Setup hugepage mappings for non-NUMA systems
[49] Setup hugepage mappings for NUMA systems
[50] Display current Ethernet/Baseband/Crypto device settings
[51] Bind Ethernet/Baseband/Crypto device to IGB UIO module
[52] Bind Ethernet/Baseband/Crypto device to VFIO module
[53] Setup VFIO permissions

-----

Step 3: Run test application for linux environment
-----
[54] Run test application ($RTE_TARGET/app/test)
[55] Run testpmd application in interactive mode ($RTE_TARGET/app/testpmd)

-----

Step 4: Other tools
-----
[56] List hugepage info from /proc/meminfo

-----

Step 5: Uninstall and system cleanup
-----
[57] Unbind devices from IGB UIO or VFIO driver
[58] Remove IGB UIO module
[59] Remove VFIO module
[60] Remove KNI module
[61] Remove hugepage mappings

[62] Exit Script

Option: 45
```

Option: 45

Unloading any existing DPDK UIO module
Loading DPDK UIO module

Press enter to continue ...

3. Bind Ethernet/Baseband/Crypto device to IGB UIO module (for this dpdk version , enter 51 to build it, Please enter the bus information of the network card that directly connected to the test machine, such as : 0000:29:00.0)

```

Option: 51

Network devices using kernel driver
=====
0000:16:00.0 'Ethernet Controller X710 for 10GBASE-T 15ff' if=ens1f0 drv=i40e unused=igb_uio *Active*
0000:16:00.1 'Ethernet Controller X710 for 10GBASE-T 15ff' if=ens1f1 drv=i40e unused=igb_uio
0000:29:00.0 'Ethernet Controller E810-C for QSFP 1592' if=ens66 drv=ice unused=igb_uio

No 'Baseband' devices detected
=====

Other Crypto devices
=====
0000:6b:00.0 'Device 4940' unused=igb_uio
0000:6d:00.0 'Device 2710' unused=igb_uio
0000:70:00.0 'Device 4940' unused=igb_uio
0000:72:00.0 'Device 2710' unused=igb_uio
0000:75:00.0 'Device 4940' unused=igb_uio
0000:77:00.0 'Device 2710' unused=igb_uio
0000:7a:00.0 'Device 4940' unused=igb_uio
0000:7c:00.0 'Device 2710' unused=igb_uio
0000:e8:00.0 'Device 4940' unused=igb_uio
0000:ea:00.0 'Device 2710' unused=igb_uio
0000:ed:00.0 'Device 4940' unused=igb_uio
0000:ef:00.0 'Device 2710' unused=igb_uio
0000:f2:00.0 'Device 4940' unused=igb_uio
0000:f4:00.0 'Device 2710' unused=igb_uio
0000:f7:00.0 'Device 4940' unused=igb_uio
0000:f9:00.0 'Device 2710' unused=igb_uio

No 'Eventdev' devices detected
=====

No 'Mempool' devices detected
=====

No 'Compress' devices detected
=====

No 'Misc (rawdev)' devices detected
=====

Enter PCI address of device to bind to IGB UIO driver: 0000:29:00.0

```

4. Setup hugepage mappings **for** NUMA systems (for this dpdk version , enter 49 to build it, Set 4G size **for** each **node**)

```

Option: 49

Removing currently reserved hugepages
Unmounting /mnt/huge and removing directory

Input the number of 1048576kB hugepages for each node
Example: to have 128MB of hugepages available per node in a 2MB huge page system,
enter '64' to reserve 64 * 2MB pages on each node
Number of pages for node0: 4
Number of pages for node1: 4
Reserving hugepages
Creating /mnt/huge and mounting as hugetlbfs

```

Modify The Configuration File

```

Enter the path where the epping program is locate:
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/apps/example/
vim epping.conf

# modify port(line 51) to your NIC PCIE number, for example:
port = 0000:29:00.0 (your NIC PCIE number)
# modify stat_print(line 105), for example:
stat_print = 0000:29:00.0

```

Configure DSA (Reconfiguration is required when restart)

You can read README in :

```
/dsa-accelerated-mtcp/dsa_userlib/README.md
```

```
# Configure dsa
# Make sure you have installed accel-config library and tools before doing this
step
cd config_dsa
./setup_dsa.sh configs/4e1w-d.conf
```

```
[root@SPR06 dsa_userlib]# cd config_dsa
[root@SPR06 config_dsa]# ./setup_dsa.sh configs/4e1w-d.conf
enabled 1 device(s) out of 1
enabled 1 wq(s) out of 1
```

Build Test Program "Epping"

```
# make test programme: epping
# epping : /mtcp_merge_virtq/app/examples/epping
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/
export RTE_SDK=`echo $PWD`/dpdk
export RTE_TARGET=x86_64-native-linuxapp-gcc
make -j32
```

Configure pktgen machine(machine run pktgen to send packets)

[Pktgen-DPDK/INSTALL.md at dev · pktgen/Pktgen-DPDK\(github.com\)](#)

1. compile DPDK:

1. Install DPDK

```
git clone https://dpdk.org/git/dpdk
sudo rm -fr /usr/local/lib/x86_64-linux-gnu # DPDK changed a number of
lib names and need to clean up
cd dpdk
meson build
ninja -C build
sudo ninja -C build install
sudo ldconfig # make sure ld.so is pointing new DPDK libraries
```

2. Compile Pktgen

```
git clone http://dpdk.org/git/apps/pktgen-dpdk
cd pktgen-dpdk
make
or
make build    # Same as 'make'
or
make rebuild  # Rebuild Pktgen, which removes the Builddir then builds
it again via meson/ninja
```

Run Demo

pktgen machine send packets ---->----->----->----->----->----->----->-----> demo spr machine
receive packets

1. pktgen machine run "pktgen-dpdk" to send packets:

```
cd /pktgen-dpdk
# start pktgen
./usr/local/bin/pktgen -c 0x3 -n 2 -- -P -m "1.0" # if the prompt "Pktgen
got a Segment Fault", try again, which may be a bug of pktgen, Until the
following interface appears
```



```

| Ports 0-0 of 1 <Main Page> Copyright(c) <2010-2021>, Intel Corporation
  Flags:Port      : P-----Sngl      :0
Link State        :      <UP-100000-FD>    ---Total Rate---
Pkts/s Rx         :      0                0
Tx                :      0                0
Mbits/s Rx/Tx     :      0/0              0/0
Pkts/s Rx Max     :      0                0
Tx Max           :      0                0
Broadcast         :      0
Multicast         :      0
Sizes 64          :      0
    65-127        :      0
    128-255       :      0
    256-511       :      0
    512-1023      :      0
    1024-1518     :      0
Runts/Jumbos      :      0/0
ARP/ICMP Pkts     :      0/0
Errors Rx/Tx      :      0/0
Total Rx Pkts     :      0
Tx Pkts           :      0
Rx/Tx MBs         :      0/0
Pattern Type      :      abcd...
Tx Count/% Rate   :      Forever /100%
Pkt Size/Tx Burst :      64 / 32
TTL/Port Src/Dest :      64/ 1234/ 5678
Pkt Type:VLAN ID  :      IPv4 / TCP:0001
802.1p CoS/DSCP/IPP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0
IP Destination    :      192.168.1.1
Source            :      192.168.0.1/24
MAC Destination   :      00:00:00:00:00:00
Source            :      b4:96:91:ad:85:b0
PCI Vendor/Addr   :      8086:1592/31:00.0
-- Pktgen 21.03.0 (DPDK 21.05.0) Powered by DPDK (pid:9275) -----

** Version: DPDK 21.05.0, Command Line Interface without timers
Pktgen:/> █

```

Enter the following command to send packets:

```

# this mac address is SPR NIC mac address, which had previously recorded
set 0 dst mac 40:A6:B7:67:19:F0
set 0 proto tcp
set 0 size 9000
start 0 rate 0.01
# if you want pause, you can enter:
stop 0

```

```

\ Ports 0-0 of 1 <Main Page> Copyright(c) <2010-2021>, Intel Corporation
  Flags:Port      : P-----Sngl      :0
Link State       :      <UP-100000-FD>    ---Total Rate---
Pkts/s Rx        :      0                0
Tx               :      8,246,017        8,246,017
Mbits/s Rx/Tx    :      0/100,271        0/100,271
Pkts/s Rx Max    :      0                0
Tx Max          :      8,246,017        8,246,017
Broadcast        :      0
Multicast        :      0
Sizes 64         :      0
65-127          :      0
128-255         :      0
256-511         :      0
512-1023        :      0
1024-1518       :      0
Runts/Jumbos     :      0/0
ARP/ICMP Pkts    :      0/0
Errors Rx/Tx     :      0/0
Total Rx Pkts    :      0
Tx Pkts          :      22,244,070
Rx/Tx MBs        :      0/270,487
Pattern Type     :      abcd...
Tx Count/% Rate  :      Forever /100%
Pkt Size/Tx Burst :      1500 / 32
TTL/Port Src/Dest :      64/ 1234/ 5678
Pkt Type:VLAN ID :      IPv4 / TCP:0001
802.1p CoS/DSCP/IPP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0
IP Destination   :      192.168.1.1
Source           :      192.168.0.1/24
MAC Destination  :      40:a6:b7:67:19:f0
Source           :      b4:96:91:ad:85:b0
PCI Vendor/Addr  :      8086:1592/31:00.0
-- Pktgen 21.03.0 (DPDK 21.05.0) Powered by DPDK (pid:9379) -----

** Version: DPDK 21.05.0, Command Line Interface without timers
Pktgen:/> set 0 dst mac 40:A6:B7:67:19:F0
Pktgen:/> set 0 proto tcp
Pktgen:/> set 0 size 1500
Pktgen:/> start 0
Pktgen:/> █

```

2. demo machine:

```

# start test produce
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/apps/example
./epping -f epping.conf

```

Then observe whether pktgen and Epping have data flow.

```

-----
Loading ARP table from : config/arp.conf
ARP Table:
IP addr: 192.168.18.124, dst_hwaddr: B4:96:91:AD:85:B1
-----
Initializing port 0... done:
[dpgk_load_module: 686] Failed to get flow control info!
[dpgk_load_module: 693] Failed to set flow control info!: errno: -95

Checking link statusdone
Port 0 Link Up - speed 100000 Mbps - full-duplex
*****DSA_INIT*****
[ info] alloc wq 0 shared 0 size 128 addr 0x7ffa2ab80000 batch sz 0x400 xfer sz 0x80000000
*****CORE INIT DSA_DONE*****
CPU 0: initialization finished.
dsa check count : 0
[mtcp_create_context:1448] CPU 0 is now the master thread.
dsa check count : 10000000
dsa check count : 20000000
dsa check count : 30000000
dsa check count : 40000000
dsa check count : 50000000
dsa check count : 60000000
TCP recvd PPS:139.91
write cnt 521402
failed cnt 513000
dsa check count : 70000000
cpu memcpy len :1048576 cost 162826 cycles
dsa check count : 80000000
write cnt 561574
failed cnt 552655
TCP recvd PPS:138.21
cpu memcpy len :1048576 cost 149142 cycles
dsa check count : 90000000
write cnt 561153
failed cnt 552240

```

4. Start Test

You must start epping first. and then start pktgen.

Sample Run

```
# Use memmove transfer size of 1m with 128 descriptors
./epping -f epping.conf -l 1m -n 128
```

The different **-n** parameter is mainly to not enter the cache when test specific transfer size(-l parameter).

benchmark with CPU

```
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/apps/example
# 4K
./epping -f epping.conf -l 4k -n 32768
# 8K
./epping -f epping.conf -l 8k -n 16384
# 16K
./epping -f epping.conf -l 16k -n 8192
# 32K
./epping -f epping.conf -l 32k -n 4096
# 64K
./epping -f epping.conf -l 64k -n 2048
# 128K
./epping -f epping.conf -l 128k -n 1024
# 256K
./epping -f epping.conf -l 256k -n 512
```

```
# 512K
./epping -f epping.conf -l 512k -n 256
# 1M
./epping -f epping.conf -l 1m -n 128
```

benchmark with DSA

```
# The performance of different parameters is different. The following is a
combination of parameters with better performance
cd /dsa-accelerated-mtcp/mtcp_merge_virtq/apps/example
# 4K
./epping -f epping.conf -l 4k -n 32768 -d
# 8K
./epping -f epping.conf -l 8k -n 16384 -d
# 16K
./epping -f epping.conf -l 16k -n 8192 -d
# 32K
./epping -f epping.conf -l 32k -n 4096 -d
# 64K
./epping -f epping.conf -l 64k -n 2048 -d
# 128K
./epping -f epping.conf -l 128k -n 1024 -d
# 256K
./epping -f epping.conf -l 256k -n 512 -d
# 512K
./epping -f epping.conf -l 512k -n 256 -d
# 1M
./epping -f epping.conf -l 1m -n 128 -d
```

5.Measurement

PPS Measurement

```
# You can observe PPS in PKTGEN
```

```

/ Ports 0-0 of 1 <Main Page> Copyright(c) <2010-2021>, Intel Corporation
Flags:Port : P-----Sngl :0
Link State : <UP-100000-FD> ---Total Rate---
Pkts/s Rx : 571,324 571,324
Tx : 160 160
Mbits/s Rx/Tx : 37,563/11 37,563/11
Pkts/s Rx Max : 583,994 583,994
Tx Max : 160 160
Broadcast : 0
Multicast : 0
Sizes 64 : 0
65-127 : 0
128-255 : 0
256-511 : 0
512-1023 : 0
1024-1518 : 0
Runts/Jumbos : 0/62,138,033
ARP/ICMP Pkts : 0/0
Errors Rx/Tx : 0/0
Total Rx Pkts : 61,567,071
Tx Pkts : 10,688
Rx/Tx MBs : 4,048,712/771
Pattern Type : abcd...
Tx Count/% Rate : Forever /0.01%
Pkt Size/Tx Burst : 9000 / 32
TTL/Port Src/Dest : 64/ 1234/ 5678
Pkt Type:VLAN ID : IPv4 / TCP:0001
802.1p CoS/DSCP/IPP : 0/ 0/ 0
VxLAN Flg/Grp/vid : 0000/ 0/ 0
IP Destination : 192.168.1.1
Source : 192.168.0.1/24
MAC Destination : 40:a6:b7:67:19:f0
Source : b4:96:91:ad:85:b0
PCI Vendor/Addr : 8086:1592/31:00.0
-- Pktgen 21.03.0 (DPDK 21.05.0) Powered by DPDK (pid:21903) -----
Pktgen:/> stop 0
Pktgen:/> start 0
Pktgen:/> stop 0
Pktgen:/> start 0
Pktgen:/> stop 0

```

CPU Usage

```

# Copy the CPU occupied by the actual operation without DSA
perf record -C 7
perf report -i perf.data

```

```

Samples: 15K of event 'cycles', Event count (approx.): 9954117485
Overhead Command Shared Object Symbol
47.87% epping libpthread-2.28.so [.] pthread_spin_lock
37.58% epping libc-2.28.so [.] __memmove_avx_unaligned_erms
3.41% epping epping [.] sfifo_out
3.34% epping epping [.] mtcp_write
3.11% epping epping [.] sfifo_in
3.07% epping [vdso] [.] __vdso_gettimeofday
0.69% epping epping [.] write_thread_func
0.31% epping epping [.] SBPut
0.10% epping libpthread-2.28.so [.] __errno_location
0.08% epping epping [.] pthread_spin_lock@plt
0.08% epping epping [.] gettimeofday@plt
0.07% epping epping [.] pthread_spin_unlock@plt
0.05% epping libpthread-2.28.so [.] pthread_spin_unlock
0.03% epping epping [.] __errno_location@plt
0.03% epping [kernel.vmlinux] [k] sync_regs
0.03% epping [kernel.vmlinux] [k] hrtimer_active
0.02% epping [kernel.vmlinux] [k] perf_event_task_tick
0.02% epping [kernel.vmlinux] [k] __intel_pmu_enable_all.constprop.48
0.01% epping [kernel.vmlinux] [k] ktime_get
0.01% epping [kernel.vmlinux] [k] __sysvec_apic_timer_interrupt
0.01% epping [kernel.vmlinux] [k] lapic_next_deadline
0.01% perf perf [.] cmd_record
0.01% epping [kernel.vmlinux] [k] rb_next
0.01% epping [kernel.vmlinux] [k] trigger_load_balance
0.01% epping [kernel.vmlinux] [k] scheduler_tick
0.01% epping [kernel.vmlinux] [k] __update_load_avg_cfs_rq
0.01% epping [kernel.vmlinux] [k] __acct_update_integrals
0.01% epping [kernel.vmlinux] [k] account_user_time
0.01% epping [kernel.vmlinux] [k] ktime_get_update_offsets_now
0.01% epping [kernel.vmlinux] [k] calc_global_load_tick
0.00% perf [kernel.vmlinux] [k] __perf_ioctl
0.00% perf [kernel.vmlinux] [k] native_apic_msr_write
0.00% perf [kernel.vmlinux] [k] smp_call_function_single

```

CPU Usage with DSA

The CPU occupied by actually initiating the DSA copy (sum of three values)

Samples: 22K of event 'cycles', Event count (approx.): 15991065519

Overhead	Command	Shared Object	Symbol
50.73%	epping	libpthread-2.28.so	[.] pthread_spin_lock
32.27%	epping	[vdso]	[.] __vdso_gettimeofday
7.08%	epping	epping	[.] write_thread_func
6.51%	epping	epping	[.] mtcp_write_async
0.93%	epping	epping	[.] sfifo_in
0.79%	epping	epping	[.] sfifo_out
0.75%	epping	epping	[.] gettimeofday@plt
0.14%	epping	libdsa_userlib.so	[.] movdir64b
0.13%	epping	[kernel.vmlinux]	[k] trigger_load_balance
0.06%	epping	epping	[.] SBPut_async
0.06%	epping	[kernel.vmlinux]	[k] ktime_get_update_offsets_now
0.05%	epping	[kernel.vmlinux]	[k] lapic_next_deadline
0.05%	epping	[kernel.vmlinux]	[k] idle_cpu
0.04%	epping	[kernel.vmlinux]	[k] ktime_get
0.03%	epping	[kernel.vmlinux]	[k] native_irq_return_iret
0.03%	epping	[kernel.vmlinux]	[k] hrtimer_active
0.03%	epping	[kernel.vmlinux]	[k] arch_scale_freq_tick
0.02%	epping	[kernel.vmlinux]	[k] send_call_function_single_ipi
0.02%	epping	[kernel.vmlinux]	[k] native_apic_msr_eoi_write
0.02%	epping	epping	[.] pthread_spin_unlock@plt
0.02%	epping	[kernel.vmlinux]	[k] sync_regs
0.01%	epping	[kernel.vmlinux]	[k] native_sched_clock
0.01%	epping	[kernel.vmlinux]	[k] _raw_spin_lock
0.01%	epping	[kernel.vmlinux]	[k] read_tsc
0.01%	epping	[kernel.vmlinux]	[k] update_irq_load_avg
0.01%	epping	[kernel.vmlinux]	[k] llist_add_batch
0.01%	epping	epping	[.] pthread_spin_lock@plt
0.01%	epping	[kernel.vmlinux]	[k] __update_load_avg_cfs_rq
0.01%	epping	[kernel.vmlinux]	[k] __sysvec_apic_timer_interrupt
0.01%	epping	[kernel.vmlinux]	[k] x2apic_send_IPI_dest
0.01%	epping	libdsa_userlib.so	[.] dsa_prep_desc_common
0.01%	epping	libdsa_userlib.so	[.] dsa_prep_memcpy
0.01%	epping	[kernel.vmlinux]	[k] swpgs_restore_regs_and_return_to_usermode
0.01%	epping	[kernel.vmlinux]	[k] __update_load_avg_se
0.01%	epping	[kernel.vmlinux]	[k] calc_global_load_tick
0.01%	epping	[kernel.vmlinux]	[k] __intel_pmu_enable_all.constprop.48
0.01%	epping	[kernel.vmlinux]	[k] _raw_spin_lock_irq
0.00%	perf	[kernel.vmlinux]	[k] __count_memcg_events.part.73
0.00%	epping	[kernel.vmlinux]	[k] hrtimer_update_next_event

Bandwidth Measurement

1. Peak value measurement

peakValue = sizeof (Total packets sent) / latency

2. Perf measurement

```
perf stat -e  
dsa0/event=0x1,event_category=0x1/,dsa0/event=0x2,event_category=0x1/ -a -I 1000
```

The final result is the counts value * 32 = 195405108 * 32