

CESG SW Hyperscale PRC Tencent Engineering Team

# How to enable IAX on SPR

Version 1.0

Ji, Yu  
8-11-2021

## How to enable IAX on SPR

### 1. Kernel build and IAX device check

CPU info:

```
Architecture:    x86_64
CPU op-mode(s):  32-bit, 64-bit
Byte Order:      Little Endian
CPU(s):          112
On-line CPU(s) list: 0-111
Thread(s) per core: 1
Core(s) per socket: 56
Socket(s):       2
NUMA node(s):    2
Vendor ID:       GenuineIntel
BIOS Vendor ID:  Intel(R) Corporation
CPU family:      6
Model:           143
Model name:      Genuine Intel(R) CPU @ 3.60GHz
BIOS Model name: Genuine Intel(R) CPU @ 3.60GHz
Stepping:        2
CPU MHz:         1700.000
CPU max MHz:     3600.0000
CPU min MHz:     800.0000
BogoMIPS:        3400.00
Virtualization:  VT-x
L1d cache:       48K
L1i cache:       32K
L2 cache:        2048K
L3 cache:        107520K
NUMA node0 CPU(s): 0-55
NUMA node1 CPU(s): 56-111
```

Flags: fpu vme de pse tsc msr pae mce cx8 apic sep mtrr pge mca cmov pat  
pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx pdpe1gb rdtscp lm  
constant\_tsc art arch\_perfmon pebs bts rep\_good nopl xtopology nonstop\_tsc cpuid  
aperfmperf tsc\_known\_freq pni pclmulqdq dtes64 monitor ds\_cpl vmx smx est tm2  
ssse3 sdbg fma cx16 xtpr pdcm pcid sse4\_1 sse4\_2 x2apic movbe popcnt  
tsc\_deadline\_timer aes xsave avx f16c rdrand lahf\_lm abm 3dnowprefetch  
cpuid\_fault epb cat\_l3 cat\_l2 cdp\_l3 invpcid\_single intel\_ppin cdp\_l2 ssbd mba ibrs  
ibpb stibp ibrs\_enhanced tpr\_shadow vnmi flexpriority ept vpid ept\_ad fsgsbase  
tsc\_adjust bmi1 avx2 smep bmi2 erms invpcid cqm rdt\_a avx512f avx512dq rdseed  
adx smap avx512ifma clflushopt clwb intel\_pt avx512cd sha\_ni avx512bw avx512vl  
xsaveopt xsavec xgetbv1 xsaves cqm\_llc cqm\_occup\_llc cqm\_mbm\_total

```
cqm_mbm_local split_lock_detect avx_vnni avx512_bf16 wbnoinvd dtherm ida arat  
pln pts hwp hwp_act_window hwp_epp hwp_pkg_req avx512vbmi umip pku ospke  
waitpkg avx512_vbmi2 gfni vaes vpclmulqdq avx512_vnni avx512_bitalg tme  
avx512_vpopcntdq la57 rdpid bus_lock_detect clDEMote movdiri movdir64b enqcmd  
fSRM avx512_vp2intersect md_clear serialize tsxldtrk pconfig arch_lbr avx512_fp16  
flush_l1d arch_capabilities
```

#### BOARD info:

```
# dmidecode 3.2  
Getting SMBIOS data from sysfs.  
SMBIOS 3.3.0 present.  
# SMBIOS implementations newer than version 3.2.0 are not  
# fully supported by this version of dmidecode.
```

Handle 0x0002, DMI type 2, 15 bytes

##### Base Board Information

```
Manufacturer: Intel  
Product Name: VulcanCity  
Version: Default string  
Serial Number: Default string  
Asset Tag: Default string  
Features:  
    Board is a hosting board  
    Board is replaceable  
Location In Chassis: Default string  
Chassis Handle: 0x0003  
Type: Motherboard  
Contained Object Handles: 0
```

Handle 0x0018, DMI type 10, 24 bytes

##### On Board Device 1 Information

```
Type: Unknown  
Status: Enabled  
Description: Device 1
```

##### On Board Device 2 Information

```
Type: Unknown  
Status: Enabled  
Description: Device 2
```

##### On Board Device 3 Information

```
Type: Unknown  
Status: Enabled  
Description: Device 3
```

##### On Board Device 4 Information

Type: Unknown  
Status: Enabled  
Description: Device 4  
On Board Device 5 Information  
Type: Unknown  
Status: Enabled  
Description: Device 5  
On Board Device 6 Information  
Type: Unknown  
Status: Enabled  
Description: Device 6  
On Board Device 7 Information  
Type: Unknown  
Status: Enabled  
Description: Device 7  
On Board Device 8 Information  
Type: Unknown  
Status: Enabled  
Description: Device 8  
On Board Device 9 Information  
Type: Unknown  
Status: Enabled  
Description: Device 9  
On Board Device 10 Information  
Type: Unknown  
Status: Enabled  
Description: <BAD INDEX>

Handle 0x0092, DMI type 41, 11 bytes

Onboard Device

Reference Designation: Onboard ETHERNET Controller

Type: Ethernet

Status: Enabled

Type Instance: 1

Bus Address: 0000:01:00.0

Handle 0x0093, DMI type 41, 11 bytes

Onboard Device

Reference Designation: Onboard VGA

Type: Video

Status: Enabled

Type Instance: 1

Bus Address: 0000:03:00.0

Handle 0x0094, DMI type 41, 11 bytes

Onboard Device

Reference Designation: Onboard SATA Controller

Type: SATA Controller

Status: Enabled

Type Instance: 1

Bus Address: 0000:00:17.0

Handle 0x0095, DMI type 41, 11 bytes

Onboard Device

Reference Designation: Onboard SATA Controller

Type: SATA Controller

Status: Enabled

Type Instance: 2

Bus Address: 0000:00:18.0

### 1.1 Configure BIOS

- "EDKII Menu" -> "Socket Configuration" -> "IIO Configuration" -> "Intel VT for directed IO (VT-d)" → Intel VT for directed IO → Enable
- "EDKII Menu" -> "Socket Configuration" -> "IIO Configuration" -> "PCI ENQCMD/ENQCMLS" → Enable
- "EDKII Menu" -> "Socket Configuration" -> "Processor Configuration" -> "Enable LP Global" → Single LP

### 1.2 Prepare to build kernel

- yum group install "Development Tools"
- yum install ncurses-devel bison flex elfutils-libelf-devel openssl-devel
- kernel version Linux 5.13.0-dsa-micros+ x86\_64
- git clone <https://github.com/torvalds/linux.git>
- git checkout v5.13
- git revert 9bfecd05833918526cc7357d55e393393440c5fa
- cd linux
- cp /boot/config-4.18.0-305.3.1.el8.x86\_64 .config
- make menuconfig
- Check build options:
  - CONFIG\_IRQ\_REMAP=y
  - CONFIG\_INTEL\_IOMMU=y
  - CONFIG\_INTEL\_IOMMU\_SVM=y
  - CONFIG\_IMS\_MSI=y
  - CONFIG\_IMS\_MSI\_ARRAY=y
  - CONFIG\_IRQ\_REMAP=y
  - CONFIG\_PCI\_ATS=y

```

CONFIG_PCI_PRI=y
CONFIG_PCI_PASID=y
CONFIG_INTEL_IDXD=m
CONFIG_INTEL_IDXD_SVM=y
CONFIG_VFIO_MDEV_IDXD=m
CONFIG_INTEL_IDXD_KTEST=m
CONFIG_DMA_ENGINE=m
CONFIG_DMATEST=m
CONFIG_IRQ_BYPASS_MANAGER=m
CONFIG_AUXILIARY_BUS=y
CONFIG_UACCE_IDXD=m
CONFIG_INTEL_IDXD_CDEV_LEGACY=y

```

- Additional options:

```

CONFIG_CRYPTODEV_IAX_CRYPTODEV=y
CONFIG_INTEL_IDXD_PAGE_CLEAR=m

```

### 1.3 Build and install kernel

- make -jX , where X is twice the number of physical cores (lscpu -> CPU(s))
- make modules
- make modules\_install
- make install
- make headers\_install INSTALL\_HDR\_PATH=/usr
- grub2-mkconfig -o /boot/grub2/grub.cfg
- grubby --set-default /boot/vmlinuz-5.13.0-dsa-micros+
- grubby --update-kernel=ALL --args="intel\_iommu=on,sm\_on"
- reboot
- uname -msr

### 1.4 IAX device check

- lspci -nn | grep 0cfe

The result should be like this

```

6a:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
6f:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
74:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
79:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
e7:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
ec:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]

```

```
f1:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
f6:02.0 System peripheral [0880]: Intel Corporation Device
[8086:0cfe]
```

You should have 2 sockets 8 IAX device

- `ls /sys/bus/dsa/devices/iax*`

The result should be like this:

```
/sys/bus/dsa/devices/iax1:
cdev_major cmd_status engine1.0 engine1.2 engine1.4
engine1.6 errors group1.0 group1.2 max_batch_size max_groups
max_transfer_size max_work_queues_size op_cap power
subsystem uevent wq1.0 wq1.2 wq1.4 wq1.6
clients configurable engine1.1 engine1.3 engine1.5 engine1.7
gen_cap group1.1 group1.3 max_engines max_tokens
max_work_queues numa_node pasid_enabled state
token_limit version wq1.1 wq1.3 wq1.5 wq1.7

/sys/bus/dsa/devices/iax11:
cdev_major cmd_status engine11.0 engine11.2 engine11.4
engine11.6 errors group11.0 group11.2 max_batch_size
max_groups max_transfer_size max_work_queues_size op_cap
power subsystem uevent wq11.0 wq11.2 wq11.4 wq11.6
clients configurable engine11.1 engine11.3 engine11.5
engine11.7 gen_cap group11.1 group11.3 max_engines
max_tokens max_work_queues numa_node pasid_enabled
state token_limit version wq11.1 wq11.3 wq11.5 wq11.7

/sys/bus/dsa/devices/iax13:
cdev_major cmd_status engine13.0 engine13.2 engine13.4
engine13.6 errors group13.0 group13.2 max_batch_size
max_groups max_transfer_size max_work_queues_size op_cap
power subsystem uevent wq13.0 wq13.2 wq13.4 wq13.6
clients configurable engine13.1 engine13.3 engine13.5
engine13.7 gen_cap group13.1 group13.3 max_engines
max_tokens max_work_queues numa_node pasid_enabled
state token_limit version wq13.1 wq13.3 wq13.5 wq13.7

/sys/bus/dsa/devices/iax15:
```

cdev\_major cmd\_status engine15.0 engine15.2 engine15.4  
engine15.6 errors group15.0 group15.2 max\_batch\_size  
max\_groups max\_transfer\_size max\_work\_queues\_size op\_cap  
power subsystem uevent wq15.0 wq15.2 wq15.4 wq15.6  
clients configurable engine15.1 engine15.3 engine15.5  
engine15.7 gen\_cap group15.1 group15.3 max\_engines  
max\_tokens max\_work\_queues numa\_node pasid\_enabled  
state token\_limit version wq15.1 wq15.3 wq15.5 wq15.7

/sys/bus/dsa/devices/iax3:

cdev\_major cmd\_status engine3.0 engine3.2 engine3.4  
engine3.6 errors group3.0 group3.2 max\_batch\_size max\_groups  
max\_transfer\_size max\_work\_queues\_size op\_cap power  
subsystem uevent wq3.0 wq3.2 wq3.4 wq3.6  
clients configurable engine3.1 engine3.3 engine3.5 engine3.7  
gen\_cap group3.1 group3.3 max\_engines max\_tokens  
max\_work\_queues numa\_node pasid\_enabled state  
token\_limit version wq3.1 wq3.3 wq3.5 wq3.7

/sys/bus/dsa/devices/iax5:

cdev\_major cmd\_status engine5.0 engine5.2 engine5.4  
engine5.6 errors group5.0 group5.2 max\_batch\_size max\_groups  
max\_transfer\_size max\_work\_queues\_size op\_cap power  
subsystem uevent wq5.0 wq5.2 wq5.4 wq5.6  
clients configurable engine5.1 engine5.3 engine5.5 engine5.7  
gen\_cap group5.1 group5.3 max\_engines max\_tokens  
max\_work\_queues numa\_node pasid\_enabled state  
token\_limit version wq5.1 wq5.3 wq5.5 wq5.7

/sys/bus/dsa/devices/iax7:

cdev\_major cmd\_status engine7.0 engine7.2 engine7.4  
engine7.6 errors group7.0 group7.2 max\_batch\_size max\_groups  
max\_transfer\_size max\_work\_queues\_size op\_cap power  
subsystem uevent wq7.0 wq7.2 wq7.4 wq7.6  
clients configurable engine7.1 engine7.3 engine7.5 engine7.7  
gen\_cap group7.1 group7.3 max\_engines max\_tokens  
max\_work\_queues numa\_node pasid\_enabled state  
token\_limit version wq7.1 wq7.3 wq7.5 wq7.7

/sys/bus/dsa/devices/iax9:

cdev\_major cmd\_status engine9.0 engine9.2 engine9.4  
engine9.6 errors group9.0 group9.2 max\_batch\_size max\_groups  
max\_transfer\_size max\_work\_queues\_size op\_cap power  
subsystem uevent wq9.0 wq9.2 wq9.4 wq9.6



|                 |              |           |             |               |            |
|-----------------|--------------|-----------|-------------|---------------|------------|
| clients         | configurable | engine9.1 | engine9.3   | engine9.5     | engine9.7  |
| gen_cap         | group9.1     | group9.3  | max_engines |               | max_tokens |
| max_work_queues | numa_node    |           |             | pasid_enabled | state      |
| token_limit     | version      | wq9.1     | wq9.3       | wq9.5         | wq9.7      |

Your IAX devices can also be in `/sys/bus/iax/devices`, it depends on the version of the kernel you use.

Useful link for more details:

<https://wiki.ith.intel.com/pages/viewpage.action?pageId=1842481266>

## 2. Enable IAX

### 2.1 Accel-config build

Repo: <https://github.com/intel/idxd-config>

- `yum groupinstall "Development Tools"`
- `yum install autoconf automake libtool pkgconf rpm-build rpmdevtools`
- `yum install asciidoc xmlto libuuid-devel json-c-devel kmod-devel libudev-devel`
- `sudo make headers_install INSTALL_HDR_PATH=/usr`
- `./autogen.sh`
- `./configure CFLAGS='-g -O2' --prefix=/usr --sysconfdir=/etc --libdir=/usr/lib64`
- `make`
- `make check`
- `sudo make install`

### 2.2 Accel-config presetting

doc for accel-config: <https://github.com/intel/idxd-config/tree/master/Documentation/accfg>

- `accel-config config-engine iax1/engine1.0 -g 0`
- `accel-config config-wq iax1/wq1.0 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y user -n app1`
- `accel-config enable-device iax1`
- `accel-config enable-wq iax1/wq1.0`

Those commands only enable one IAX device(iax1) and one work queue(wq1.0) on one engine(engine1.0).

If you want all 8 iax devices enabled with 8 engines and 8 work queues, you need to do this in the shell:

```
for ((i=1;$i<=15;i=$i+2))
do
    echo $i

    accel-config config-engine iax$i/engine$i.0 -g 0
    accel-config config-engine iax$i/engine$i.1 -g 0
    accel-config config-engine iax$i/engine$i.2 -g 0
    accel-config config-engine iax$i/engine$i.3 -g 0
    accel-config config-engine iax$i/engine$i.4 -g 0
    accel-config config-engine iax$i/engine$i.5 -g 0
    accel-config config-engine iax$i/engine$i.6 -g 0
    accel-config config-engine iax$i/engine$i.7 -g 0

    accel-config config-wq iax$i/wq$i.0 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.1 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.2 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.3 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.4 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.5 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.6 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1
    accel-config config-wq iax$i/wq$i.7 -g 0 -s 16 -p 10 -b 1 -t 15 -m shared -y
    user -n app1

    accel-config enable-device iax$i

    accel-config enable-wq iax$i/wq$i.0
    accel-config enable-wq iax$i/wq$i.1
    accel-config enable-wq iax$i/wq$i.2
    accel-config enable-wq iax$i/wq$i.3
    accel-config enable-wq iax$i/wq$i.4
```

```
accel-config enable-wq iax$/wq$.5
accel-config enable-wq iax$/wq$.6
accel-config enable-wq iax$/wq$.7
done
```

Check if the IAX devices are enabled by:

- accel-config list

### 3. QPL build and test

Register the access through:

<https://registrationcenter.intel.com/en/forms/?productid=3389&pass=yes>

Password: **accept**

Please be note, QPL is in alpha phase and not publicly release, persons in below list will get approved to access QPL.

```
ryanxgao@tencent.com ruanxgao
frankpang@tencent.com frank
felixtwang@tencent.com felix
hardyffu@tencent.com hardy
ivanren@tencent.com ivanren
ruippan@tencent.com ruip
```

#### 3.1 Code modification

The test code is in ***qpl\_library/examples/high-level-api/simple-operations***. Find the function you want to test and replace ***software*** as ***hardware*** in the code.

Take ***compression\_example.cpp*** as an example, your code after modification should be like this:

```
* were provided to you ("License"). Unless the License provides otherwise,
* you may not use, modify, copy, publish, distribute, disclose or transmit
* this software or the related documents without Intel's prior written
* permission.
*
* This software and the related documents are provided as is, with no
* express or implied warranties, other than those that are expressly
* stated in the License.
*
*/

/** [QPL_HIGH_LEVEL_COMPRESSION_EXAMPLE] */

#include <iostream>
```

```

#include <string>

#include <operations/compression/deflate_operation.hpp>
#include <operations/compression/inflate_operation.hpp>

auto main() -> int {
    // Source and output containers
    std::vector<uint8_t> source(1000, 5);
    std::vector<uint8_t> destination(500, 4);
    std::vector<uint8_t> reference(1000, 7);

    // Performing an operations
    auto deflate_operation = qpl::deflate_operation::builder()
        .compression_level(qpl::compression_levels::default_level)
        .compression_mode<qpl::compression_modes::dynamic_mode>()
        .gzip_mode(false)
        .build();
    auto inflate_operation = qpl::inflate_operation();

    const auto compressed_result =
qpl::execute<qpl::hardware>(deflate_operation,
                            std::begin(source),
                            std::end(source),
                            std::begin(destination),
                            std::end(destination));

    // Handle compression result
    compressed_result.handle([](uint32_t value) -> void {
        std::cout << "Compressed size: " << value << std::endl;
    },
    [](uint32_t status_code) -> void {
        throw std::runtime_error("Error: Status code - " +
std::to_string(status_code));
    });

    const auto decompressed_result =
qpl::execute<qpl::hardware>(inflate_operation, destination, reference);

    // Handle decompression result
    decompressed_result.handle([&source, &reference](uint32_t
decompressed_size) -> void {
        // Check if everything was alright
        if (decompressed_size != source.size()) {

```

```

        throw std::runtime_error("Decompressed buffer size is
not the same "
                                "as source size.");
    }

    // Compare source and decompressed buffers
    for (size_t i = 0; i < source.size(); i++) {
        if (source[i] != reference[i]) {
            throw std::runtime_error("Content wasn't successfully
compressed "
                                    "and decompressed.");
        }
    }
},
[](uint32_t status_code) -> void {
    throw std::runtime_error("Error: Status code - " +
std::to_string(status_code));
});

    std::cout << "Content was successfully compressed and decompressed." <<
std::endl;

    return 0;
}

/* [QPL_HIGH_LEVEL_COMPRESSION_EXAMPLE] */

```

Notice that *qpl::software* have been replaced by *qpl::hardware* .

### 3.2 Build and test

- mkdir build
- cd build
- cmake -DCMAKE\_INSTALL\_PREFIX=/usr/local ..
- cmake --build . --target install
- cd examples/high-level-api

Now you can run the test you want. The result should be like this:

```

[root@spr02 high-level-api]# ./compression_example
Compressed size: 33
Content was successfully compressed and decompressed.

```